



**Budapest University of Technology and Economics**  
Faculty of Electrical Engineering and Informatics  
Department of Networked Systems and Services

# Automated monitoring of IT systems and network devices

BACHELOR'S THESIS

*Author*

Miklós Kozák

*Advisor*

Dr. Zoltán Zsóka

December 13, 2022

# Contents

<b>Kivonat</b>	<b>i</b>
<b>Abstract</b>	<b>ii</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Monitoring Systems . . . . .	1
1.2 Computer Engineering Special Interest Group of Schönherz Zoltán Dormitory	1
1.2.1 Infrastructure . . . . .	1
1.2.2 AuthSCH: Authentication and Authorization . . . . .	2
1.2.3 osTicket: Support Ticketing System . . . . .	2
1.2.4 Mattermost: Main Communication Platform . . . . .	2
1.2.5 GitLab: DevOps Platform . . . . .	2
1.3 Overview of this Thesis . . . . .	2
<b>2 Earlier Monitoring Systems Implemented in the Dormitory</b>	<b>3</b>
2.1 Nagios / Icinga (2009(?) - 2015) . . . . .	3
2.1.1 Reasons for Phasing Out . . . . .	3
2.1.2 Compared to Modern Standards . . . . .	4
2.2 Zabbix (2015 - 2019) . . . . .	4
2.2.1 Reasons for Phasing Out . . . . .	4
2.2.2 Compared to Modern Standards . . . . .	4
2.3 Prometheus (2019 - 2021) . . . . .	4
2.3.1 Reasons for Phasing Out . . . . .	5
2.3.2 Compared to Modern Standards . . . . .	5
2.4 Lessons learned . . . . .	5
<b>3 Requirements of the New Monitoring System</b>	<b>6</b>
3.1 Collecting and Storing Metrics and Logs . . . . .	6
3.2 Effective long-term storage of historical data . . . . .	7
3.3 Detecting issues and alerting the necessary personnel . . . . .	7

3.4	Detecting issues of the monitoring system itself (meta-monitoring) . . . . .	7
3.5	Ease of use yet highly customizable . . . . .	8
3.6	Authentication and authorization via AuthSCH . . . . .	8
3.7	Visualization of collected data . . . . .	8
3.8	Security . . . . .	8
<b>4</b>	<b>Available industry standard technologies</b>	<b>10</b>
4.1	Prometheus vs VictoriaMetrics . . . . .	10
4.1.1	Prometheus . . . . .	10
4.1.2	VictoriaMetrics . . . . .	10
4.1.3	Comparison . . . . .	10
4.1.4	Verdict . . . . .	11
4.2	ELK stack vs Loki and Promtail . . . . .	11
4.2.1	ELK stack . . . . .	11
4.2.2	Loki and Promtail . . . . .	11
4.2.3	Comparison . . . . .	11
4.2.4	Verdict . . . . .	11
4.3	Grafana . . . . .	12
<b>5</b>	<b>System architecture</b>	<b>13</b>
5.1	System Components . . . . .	13
5.1.1	VictoriaMetrics Components . . . . .	13
5.1.1.1	VMAgent . . . . .	13
5.1.1.2	VMSingle . . . . .	14
5.1.1.3	VMAlert . . . . .	14
5.1.1.4	VMAlertmanager . . . . .	14
5.1.2	Promtail . . . . .	14
5.1.3	Loki . . . . .	14
5.1.4	Grafana . . . . .	14
5.2	Deployment . . . . .	14
5.2.1	Kubernetes . . . . .	15
<b>6</b>	<b>Implementation details and experiences</b>	<b>16</b>
6.1	The server hardware and the host system . . . . .	16
6.1.1	The Hardware . . . . .	16
6.1.2	The Operating System . . . . .	16
6.1.3	The Kubernetes Distribution . . . . .	17
6.1.4	The Final Host System Configuration . . . . .	17

6.2	Preparing the k3s Kubernetes distribution . . . . .	17
6.2.1	Management Tools . . . . .	17
6.2.2	Persistent Volumes . . . . .	18
6.3	Installing and configuring a basic VictoriaMetrics on k3s . . . . .	18
6.3.1	Downloading the Helm Charts . . . . .	19
6.3.2	Installing VictoriaMetrics . . . . .	19
6.4	Collecting metrics . . . . .	19
6.4.1	Collecting Metrics from Servers . . . . .	19
6.4.2	Collecting Metrics from Network Devices . . . . .	20
6.4.3	Metric Visualizations . . . . .	22
6.5	Automating tasks with the help of GitLab CI . . . . .	22
6.5.1	Implementing Helm Charts for Configuration Generation . . . . .	23
6.5.2	Configuring GitLab CI . . . . .	24
6.6	Authentication and authorization . . . . .	26
6.7	Collecting logs . . . . .	26
6.7.1	Configuring and Installing Loki . . . . .	26
6.7.2	Collecting Logs of the Monitoring System . . . . .	28
6.7.3	Receiving and Processing Logs from other Systems . . . . .	28
6.7.4	Log Visualization . . . . .	31
6.7.5	Security . . . . .	31
6.8	Creating alerts and sending notifications . . . . .	32
6.8.1	Configuring VMAAlert . . . . .	32
6.8.2	Configuring VMAAlertmanager . . . . .	33
6.8.3	Alerts on Mattermost . . . . .	35
6.8.4	Implementing Continuous Integration for Alerts . . . . .	35
<b>7</b>	<b>Results and evaluation</b>	<b>37</b>
7.1	Evaluation of Requirements Fulfillment . . . . .	37
7.1.1	Collecting and storing metrics and logs . . . . .	37
7.1.2	Effective long-term storage of historical data . . . . .	37
7.1.3	Detecting issues and alerting the necessary personnel . . . . .	37
7.1.4	Detecting issues of the monitoring system itself (meta-monitoring) . . . . .	38
7.1.5	Ease of use yet highly customizable . . . . .	38
7.1.6	Authentication and authorization via AuthSCH . . . . .	38
7.1.7	Visualization of collected data . . . . .	38
7.1.8	Security . . . . .	38
7.1.9	Summary . . . . .	38

7.2	Evaluation of the Performance of the System . . . . .	38
7.2.1	CPU and RAM Usage . . . . .	38
7.2.2	Disk Usage . . . . .	39
7.2.3	Network Usage . . . . .	39
7.2.4	Summary . . . . .	39
7.3	Results of Continuous Integration . . . . .	39
7.3.1	CI for Collecting Metrics from Servers . . . . .	39
7.3.2	CI for Collecting Metrics from Network Devices . . . . .	40
7.3.3	CI for Configuring Alerts . . . . .	40
7.3.4	Summary . . . . .	40
7.4	Other Results of the Monitoring System running so far . . . . .	40
<b>8</b>	<b>Future Development Opportunities</b>	<b>41</b>
8.1	Improving Configuration Management and Enforcing Reproducibility . . . .	41
8.1.1	Grafana Configurations . . . . .	41
8.1.2	Backups . . . . .	41
8.1.3	Ansible Scripts . . . . .	41
8.1.4	Advanced Configuration Generation . . . . .	42
8.2	Granting Users Access to Internal Components . . . . .	42
8.2.1	Access to the VMAgent component . . . . .	42
8.2.2	Access to the VMAlert component . . . . .	42
8.2.3	Access to the VMAlertmanager component . . . . .	42
8.3	Improving Log Collection . . . . .	42
8.3.1	Configurations for Remote Promtails . . . . .	42
8.3.2	Collecting Logs from Network Devices . . . . .	43
8.4	Advanced Log Analysis . . . . .	43
8.5	Service Discovery . . . . .	43
8.6	Other Possibilities . . . . .	43
<b>9</b>	<b>Conclusion</b>	<b>44</b>
	<b>Acknowledgements</b>	<b>45</b>
	<b>Bibliography</b>	<b>46</b>
	<b>Appendix</b>	<b>47</b>
A.1	Initial VictoriaMetrics values.yml . . . . .	47
A.2	SNMP scraper helm chart . . . . .	50

A.3 Alerts helm chart . . . . .	52
---------------------------------	----

## HALLGATÓI NYILATKOZAT

Alulírott *Kozák Miklós*, szigorló hallgató kijelentem, hogy ezt a szakdolgozatot meg nem engedett segítség nélkül, saját magam készítettem, csak a megadott forrásokat (szakirodalom, eszközök stb.) használtam fel. Minden olyan részt, melyet szó szerint, vagy azonos értelemben, de átfogalmazva más forrásból átvettem, egyértelműen, a forrás megadásával megjelöltem.

Hozzájárulok, hogy a jelen munkám alapadatait (szerző(k), cím, angol és magyar nyelvű tartalmi kivonat, készítés éve, konzulens(ek) neve) a BME VIK nyilvánosan hozzáférhető elektronikus formában, a munka teljes szövegét pedig az egyetem belső hálózatán keresztül (vagy autentikált felhasználók számára) közzétegye. Kijelentem, hogy a benyújtott munka és annak elektronikus verziója megegyezik. Dékáni engedéllyel titkosított diplomatervek esetén a dolgozat szövege csak 3 év eltelte után válik hozzáférhetővé.

Budapest, 2022. december 13.

---

*Kozák Miklós*  
hallgató

# Kivonat

Az automatikus felügyelő rendszerek feladata metrikákat gyűjteni hardverekről és szolgáltatásokról, hibákat találni és megpróbálni előre megjósolni a jövőben bekövetkező meghibásodásokat. Amennyiben szükséges értesítik a megfelelő személyzetet, hogy gyorsan intézkedhessenek ezáltal megakadályozva, hogy a probléma túl sok felhasználóra legyen hatással. A modern rendszerek támogatják továbbá a naplóbejegyzések gyűjtését és fejlett elemzését, illetve minden elérhető adat vizualizációját.

A Kollégiumi Számítástechnikai Kör (röviden KSZK) aktív tagsága nagyjából 20 fő. Ők a Villamosmérnöki és Informatikai Kar olyan hallgatói, akik önkéntesként teszik munkájukat. Ők különféle informatikai szolgáltatásokat üzemeltetnek a VIK hallgatói számára és emellett a Schönherz Zoltán Kollégium teljes hálózatát is üzemeltetik.

A szakdolgozatomban egy automatikus felügyelő rendszer megvalósítását mutatom be. Az elkészült rendszerben VictoriaMetrics gyűjti a metrikákat a szerverekről és a hálózati eszközökről egyaránt. A VMAalertmanager komponens küldi a riasztási értesítéseket egy Mattermost szerverre. Loki és Promtail gyűjti be és dolgozza fel a naplóbejegyzéseket. Folyamatos Integrációt használok új felügyelési célpontok bekötésére a futó rendszerbe úgy, hogy egy felügyelő rendszer adminisztrátor közbenjárására sincs szükség.

Ez az új rendszer remélhetőleg segíteni fog a KSZK-nak emelni szolgáltatásai színvonalát miközben csökkenti a tagjaira hárult terheket.

Ebben a dolgozatban leírom a teljes folyamatot, melynek eredménye a fent leírt rendszer implementálása. Tanulmányozom a KSZK által korábban megvalósított rendszereket annak érdekében, hogy rájöjjek mi kell egy sikeres megvalósításhoz. Ezen eredmények alapján leírom a KSZK-ban futó felügyelő rendszer kemény követelményeit és felkutatok olyan modern technológiákat, amelyek ki tudják elégíteni ezen követelményeket. Az így szerzett tudást felhasználva meghatározom egy VictoriaMetrics alapú rendszer szerkezetét, majd elkezdem kidolgozni a részleteket és megvalósítani a tervet. Értékelem az új rendszert és ellenőrzöm az elért eredményeket. Végül kifejtem miket lehetne még tenni a jövőben, beleértve lehetséges továbbfejlesztéseket és teljesen új ötleteket is.

A dolgozat során elkészült rendszer megfelel a felállított követelményeknek. A KSZK tagjai már használatba is vették az elkészült funkciókat és a Folyamatos Integráció segítségével sok új rendszert kezdtek el megfigyelni.



# Abstract

Automated monitoring systems collect metrics from hardwares and services, detect error conditions and try to predict future failures. If necessary the right personnel are notified to take immediate action preventing the issue from affecting too many users. Modern systems also support the collection of logs as well as advanced analysis and visualization of every available data.

Computer Engineering Special Interest Group (CESIG for short) is a team of around 20 day-to-day active members. They are volunteers from the students of Faculty of Electrical Engineering and Informatics. They provide various IT services for the Faculty's students and are responsible for the operation of Schönherz Zoltán Dormitory's whole network.

In this thesis I implement an automated monitoring system. VictoriaMetrics scrapes metrics from servers and network devices. VMAalertmanager sends alert notifications to a Mattermost server. Loki and Promtail collect and process logs. Grafana visualizes collected metrics and logs. Continuous Integration is used to integrate new monitoring targets into the running system without the intervention of a monitoring system administrator.

This new system will hopefully help CESIG improve the quality of their services while reducing workload on members.

In this thesis I describe the whole process resulting in the implementation of the system described above. I study earlier takes on implementing such systems by CESIG to figure out what is needed for a successful implementation. Based on these results I establish the hard requirements for a monitoring system in CESIG's infrastructure and explore modern technologies that can satisfy these requirements. Using the knowledge gained I describe the new system architecture for a VictoriaMetrics based solution and implement the specified architecture and work out the details. I evaluate the new system and check out the results. Finally I future work including possible improvements and new ideas as well.

The implemented system satisfies the requirements established in this thesis. Members of CESIG are already using the completed functions and started monitoring various new systems with Continuous Integration.

# Chapter 1

## Introduction

### 1.1 Monitoring Systems

Maintaining one or two servers for personal purposes is not a daunting task. We can easily keep an eye on our hardwares and apply fixes when necessary. And if for some reason a server goes down it is not likely to cause issues for others.

In a real-world production infrastructure the situation is radically different. Manually checking each physical machine's state is an overwhelming burden and it quickly becomes unfeasible. Sooner or later problems that could have been prevented are going to cause service disruptions. This is going to cause issues for users and the maintainers will only be notified through users' complaints. At this point in time it is too late as the service's reputation is damaged and the users are dissatisfied.

The best solution to this is an automated monitoring system. These software stacks collect metrics from servers and other devices, detect error conditions and try to predict future failures. If necessary, the personnel in charge are notified to take immediate action preventing the issue from affecting too many users. Modern systems also support the collection of logs as well as advanced analysis and visualization of all available data.

### 1.2 Computer Engineering Special Interest Group of Schönherz Zoltán Dormitory

Computer Engineering Special Interest Group (CESIG for short) is a team of around 20 day-to-day active members. They are volunteers from the students of Faculty of Electrical Engineering and Informatics. They provide various IT services for the Faculty's students and are responsible for the operation of Schönherz Zoltán Dormitory's whole network.

The monitoring system designed and implemented in this thesis will hopefully help CESIG improve the quality of their services while reducing workload on members.

#### 1.2.1 Infrastructure

The IT infrastructure is comparable to that of a medium sized office building. CESIG's Sysadmin group operates around 20 physical servers running virtualization clusters, storage clusters and much more. CESIG's NETeam group operates around 10 physical servers

and 50 network devices serving more than 1000 endpoints for students living inside the dormitory .

### **1.2.2 AuthSCH: Authentication and Authorization**

AuthSCH is the central authentication and authorization service. It was created by Zoltán Janega [4], a veteran member of CESIG to replace the 3 authentication systems used back then. Users data is stored in Active Directory and authentication is done over the OAuth 2.0 protocol.

### **1.2.3 osTicket: Support Ticketing System**

Users' issues are reported through an osTicket instance. Between 20 to 100 tickets are opened in each month, depending on which part of a semester we are in. Answering and resolving these tickets take a lot of time.

### **1.2.4 Mattermost: Main Communication Platform**

Mattermost is an open-source communication platform with a rich feature set. Communications between the members of CESIG mainly take place on a self-hosted Mattermost server besides e-mails.

### **1.2.5 GitLab: DevOps Platform**

GitLab is an open-core DevOps platform. It's not only a git server as it features an issue tracker, Continous Integration and much more. CESIG operates a self-hosted GitLab server for storing their source codes and configurations and to run CI jobs and more.

## **1.3 Overview of this Thesis**

Monitoring systems are not unknown to CESIG . In Chapter 2 we are going to study earlier takes on implementing such systems in the dormitory to figure out what is needed for a successful implementation.

In Chapter 3 we are going to define the hard requirements for a monitoring system in CESIG 's infrastructure.

After that we are going to explore modern technologies suitable for our needs in Chapter 4.

Using the results of our research in previous chapters we are going to describe the new system architecture in Chapter 5.

In Chapter 6 we are going to start implementing the specified architecture and work out the details.

We are going to evaluate the new system in Chapter 7 and check out the results.

Finally, in Chapter 8 we are going to discuss future work including possible improvements and new ideas as well.

## Chapter 2

# Earlier Monitoring Systems Implemented in the Dormitory

The demand for monitoring of live services rose in Computer Engineering Special Interest Group just like it rose in the IT industry outside the Schönherz Zoltán Dormitory . To satisfy this demand CESIG 's members started to operate such systems. We are going to examine the most recent of these systems and take a look at why they were phased out.

Note that these systems had been used long ago. These softwares changed a lot since then. We are going to asses them as they were seen by CESIG 's member back then. The following sections are based on archived internal notes, e-mail threads and the memories of veteran CESIG members.

### 2.1 Nagios / Icinga (2009(?) - 2015)

The oldest system we have information about was a Nagios based monitoring appliance. It was a monolithic service written in C. Later on, Nagios was changed for it's open source fork, Icinga. The probable reasons for this were a better User Interface and better User Experience overall.

Nagios/Icinga served as the monitoring and alerting component of the stack while plain old syslog-ng collected logs from remote hosts.

#### 2.1.1 Reasons for Phasing Out

Many services were not connected to this system and this was a recurring issue back then.

Another problem was that every alert and notification flooded a single mailing list. Every user was expected to filter these mails themselves. The high level of noise made real alerts hard to see and thus they were often ignored.

It's web interface was also very primitive. It did not support authorization, customization nor visualization by itself.

### **2.1.2 Compared to Modern Standards**

Comparison is difficult as this was set up much more than 10 years ago. System administration was very different back then, long before Docker, Ansible and Continuous Integration.

## **2.2 Zabbix (2015 - 2019)**

Icinga's successor system was Zabbix based. Zabbix heavily relied on Java back then and was also monolithic. It was a much more advanced system than Icinga, featuring built-in visualization, customizable user interface and authorization. This system sent alerts directly to system administrators. The advantage of this approach was less noise while the downside was less transparency.

Log collecting was done the same way as before, with no advanced processing or querying. Only a forgotten Perl script made some daily statistics.

### **2.2.1 Reasons for Phasing Out**

A huge downside of this system was its high system requirements. Large amounts of I/O made SSDs mandatory which wore out quickly. Not only I/O but RAM and CPU usage were intensive as well.

Another downside was the Zabbix monitoring daemons. These were heavyweight Java applications shipped with their own JVMs. These pushed metrics to Zabbix.

Many systems were not connected to this system either.

### **2.2.2 Compared to Modern Standards**

Compared to newer systems several other downsides became apparent.

Configurations were stored internally and thus configuration management would have been very hard if not impossible. The push-style metric collecting also made service discovery difficult. The monolithic design and the usage of a relational database made scaling very difficult compared to newer stacks using clusters of smaller services and time series databases.

On the other hand its web UI was similar to modern visualization tools featuring user configurable dashboards and control panels.

## **2.3 Prometheus (2019 - 2021)**

The most recent system was Prometheus (Section 4.1.1) based with an Elasticsearch-Logstash-Kibana stack (Section 4.2.1) for receiving, processing and querying logs and Grafana (Section 4.3) for visualization, running on plain Docker. A standard Prometheus monitoring system architecture.

The system featured Configuration Management and Continuous Integration. These together ensured easier maintenance and usage of the system. Every tool and configuration

was stored in multiple repositories on CESIG 's GitLab server. CI and advanced configuration generation made the integration of new systems much easier.

It was capable of monitoring Linux and Windows servers, network devices and even some very exotic devices of the dormitory .

Alert notifications were sent to relevant channels on CESIG 's Mattermost as well as to other communications channels as required. This made alerts transparent and much less noisy.

Higher number of services were integrated into this system than ever before, thanks to the above properties.

### **2.3.1 Reasons for Phasing Out**

In order to achieve it's great features the system made use of many unique, customly created components. These included Docker container management, configuration generation and others. The downside of this solution was a very steep learning curve. Maintenance of this stack required the understanding of these tools. This would take too many hours for newcomers.

For this reason and miscommunications the system was demolished but a new system was never made to take it's place until now. Luckily this time every tool and configuration is available to us in the project's git repositories.

### **2.3.2 Compared to Modern Standards**

This software stack had all the attributes expected from a modern system. Authentication and authorization, a time series database, smaller services, scalability and more.

## **2.4 Lessons learned**

If we want a monitoring system to be fully utilized and maintained for a long period of time it needs to have some key properties. These are:

- A way to add new systems with minimal effort
- Lightweight, simple tools for monitoring targets
- Industry standard tools instead of custom ones to achieve a shallow learning curve
- Coherent documentation
- Useful alert notifications targeted only at those who are interested in it
- Interactive, easy to use and highly customizable web UI
- Pretty visualization

## Chapter 3

# Requirements of the New Monitoring System

Based on lessons learned from previous implementations and keeping in mind the structure and aims of the Computer Engineering Special Interest Group we can define the hard and soft requirements of a new and effective monitoring system.

### 3.1 Collecting and Storing Metrics and Logs

First of all, we have to collect various kinds of metrics logs from many kinds of services and hardware devices. This is essential in achieving the main aims of this project.

Hardware devices include but are not limited to:

- servers running Windows Operating Systems
- servers running Linux Operating Systems
- servers running VMWare ESXi
- various Cisco switches, routers and access points
- other exotic hardwares such as the ones observing the washing machines and dryers of the dormitory

Services include but are not limited to:

- various databases: MySQL, PostgreSQL, etc.
- virtualization clusters: VMWare, Hyper-V, etc.
- the Kubernetes cluster
- storage servers and storage clusters
- the mailing service and mailing lists
- the dormitory's network and internet access
- unique services developed for internal use

Collecting metrics from higher level devices such as Linux hosts can be achieved in many ways as these can run almost anything. The most common way of achieving this is by running one or more monitoring daemons on the specific host. These daemons expose the metrics via an HTTP interface from which the monitoring solution pulls the metrics from. As for network devices the options are limited as often it is impossible to run arbitrary software on them. The industry standard way of probing these devices is via the SNMP protocol.

Collecting logs on the other hand is the other way around: devices push their logs onto the centralized log collector server. Servers and such have many options to chose from just like before. The industry standard solution for network devices is via the syslog protocol.

Collecting metrics and logs (if applicable) from unique devices and services might require special solutions and thus the system must be highly flexible.

### **3.2 Effective long-term storage of historical data**

Even though the dormitory's IT systems are small compared to a larger IT company the sheer amount of metrics and logs collected from every device and service can reach huge sizes on the long run. We must keep in mind that the CESIG 's resources are limited.

In order to minize the storage requirements the system must be capable of selective retention of the collected data. Per-second CPU metrics don't hold much value after a day or so while firewall logs might be valuable even after years.

Older data is much less frequently accessed. The ability to store these entries highly compressed and on another dedicated storage server would be much more cost-effective and thus highly desired.

### **3.3 Detecting issues and alerting the necessary personnel**

Through analyzation of the collected data we can identify coming and ongoing issues. Automatic detection of these problems and the alerting of the right CESIG members will prevent failures, improve the uptime and quality of our services and thus keeping our users more happy.

Care must be taken though. False alerts or notifications sent to the wrong maintainers will decrease confidence in the system or worse real malfunctions are going to be ignored.

The CESIG 's main internal communication platforms are via e-mail and Mattermost. The system must be able to use these platforms at least.

### **3.4 Detecting issues of the monitoring system itself (meta-monitoring)**

Relying on automated monitoring can make the life of CESIG 's network and system administrators much easier but this dependency has a major risk. In case parts or whole of the monitoring system fails no automatic warnings will be generated. Another issue is that the monitoring server is located in the CESIG 's server room inside the Schönherz Zoltán Dormitory . In case the dormitory 's internet connection is severed the system can not send out alerts.



Monitoring of the monitoring system itself (either by the system monitoring itself or two systems monitoring each other) is called meta-monitoring.

Partial problems of the system (e.g. one effecting only the log collector) can likely be detected by the system monitoring itself. On the other hand a critical failure can only be detected by a separate, much simpler system running offsite.

### **3.5 Ease of use yet highly customizable**

Not only the CESIG 's hardware resources are limited so are it's members time. For this reason the configuration of metric scraping, log collection and alerting must be easy and should not require any special knowledge from anybody other than the monitoring system's administrator(s).

At the same time the system must be highly customizable to accomodate unique solutions.

This can be achieved by using industry-standard solutions (eliminating a steep learning curve) and fully automating common configurations while leaving custom solutions possible.

### **3.6 Authentication and authorization via AuthSCH**

The system thus must be able to authenticate and authorize it's users via OAuth 2.0.

Based on the user's AD groups:

- monitoring system admins must have superadmin privileges
- CESIG full members must be able to view everything and modify existing services or add new ones
- CESIG fresh recruits must be able to view every log and metric
- everyone else must be denied access

### **3.7 Visualization of collected data**

All the collected data would be meaningless without proper interpretation and visualization. Thus it is crucial for the system to provide a modern, web-based and easy to use interface for visualizing collected metrics and browsing logs. Users should be able to create new dashboards, access logs and much more without the help of the monitoring admins.

Another important aspect of a pretty and modern interface are public relations. With easy to grasp graphs, diagrams and numbers it's much simpler to explain the countless hours the CESIG 's voluntary administrators put into it's many services. More importantly it could inspire the upcoming generation to take part in our work and learn with us.

### **3.8 Security**

As this system must have access to many internal systems and resources proper security is crucial. Not only the system must be defended from outside harm, it must also validate

the authenticity of every metric and log it has access to. Without proper measures an adversary could inject false information into the system or much worse gain access to sensitive information.

## Chapter 4

# Available industry standard technologies

In this chapter we are going to view and compare some technologies that are able to satisfy our requirements.

### 4.1 Prometheus vs VictoriaMetrics

#### 4.1.1 Prometheus

Prometheus[6] is an open-source monitoring system that is designed to monitor the performance and health of a computer system. It is commonly used in large, complex environments and is known for its scalability, reliability, and ability to provide real-time alerts and notifications when problems arise.

#### 4.1.2 VictoriaMetrics

VictoriaMetrics[8] is an open-source time-series database and a monitoring solution that is designed to be highly efficient and scalable. Additionally, it offers a number of features that make it easy to manage and maintain, such as automatic data retention and garbage collection.

#### 4.1.3 Comparison

Both systems are capable of performing the needed tasks. As VictoriaMetrics uses and/or compatible with Prometheus APIs it can work with the exact same software environments and can be a drop-in replacement. This is handy for us as CESIG's system administrators are already familiar with the needed exporters.

VictoriaMetrics has some clear advantages over Prometheus for us[5]:

- it's local storage can work as a durable long-term storage
- it utilizes CPU and RAM better

#### 4.1.4 Verdict

Although CESIG had utilized Prometheus in the past we decided to use VictoriaMetrics for this project.

## 4.2 ELK stack vs Loki and Promtail

### 4.2.1 ELK stack

An ELK stack[1] is a term used to refer to a combination of three open-source tools: Elasticsearch, Logstash, and Kibana. These tools are used for different purposes, but they are often used together to create a powerful logging and monitoring solution. Elasticsearch is a search engine and data analysis tool that is used to store and index large volumes of data. Logstash is a data collection and processing tool that is used to ingest, transform, and filter data from a variety of sources. Kibana is a data visualization tool that is used to create dashboards and other visualizations from the data stored in Elasticsearch.

### 4.2.2 Loki and Promtail

Loki and Promtail[3] are part of a set of open-source tools for log management. Loki is a horizontally-scalable, highly-available log aggregation system. It is designed to be very efficient and can handle large amounts of log data. Promtail is a log scraping agent that can collect log data from a variety of sources, including containers and Kubernetes environments. It is designed to work seamlessly with Loki, allowing you to easily ingest and query your log data.

### 4.2.3 Comparison

ELK and Loki serve the same purpose yet they are completely different.

Advantages of Loki over ELK:

- easy to operate
- requires much less resources
- collected logs can be visualized by Grafana
- authorization (through Grafana)

Advantages of ELK over Loki:

- supports much more complex searches

### 4.2.4 Verdict

For our use cases Loki's capabilities seem fine although we can not see the future. On the other hand an ELK stack requires huge amounts of resources which would be a problem soon. We decided to use Loki.

## 4.3 Grafana

Grafana[2] is an open-source data visualization tool known for its powerful and customizable dashboard features, which allow users to create real-time graphs, charts, and other visualizations from the data collected by other tools.

Grafana supports a large array of data sources including Prometheus, VictoriaMetrics and Loki. It also supports authentication and authorization through OAuth 2.0.

## Chapter 5

# System architecture

As discussed in the previous chapter we have decided to use VictoriaMetrics, Loki, Promtail and Grafana for our new monitoring system. Let's see the new system architecture.

### 5.1 System Components

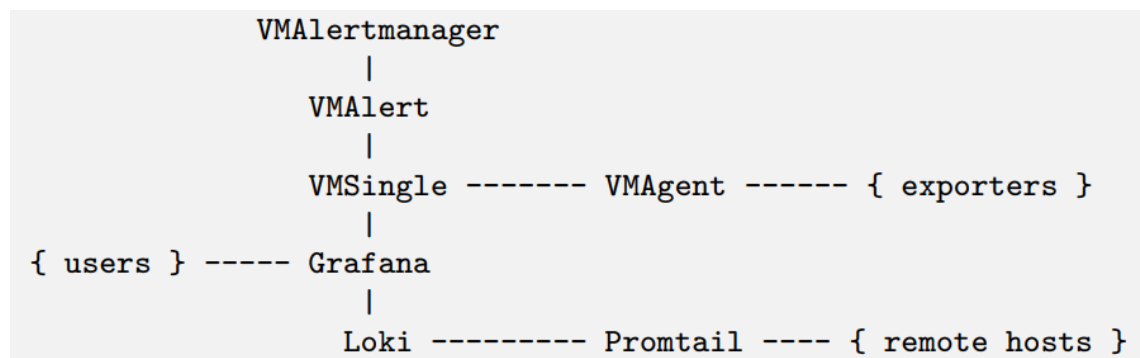


Figure 5.1: System Components

#### 5.1.1 VictoriaMetrics Components

We are going to use VictoriaMetrics for scraping and storing metrics, generating alerts and sending out notifications.

As CESIG's infrastructure is small compared to what VictoriaMetrics is able to handle we have opted for a simple deployment. This type of deployment uses one service of each type which is easier to configure and handle at the cost of scalability. We do not expect to hit any bottlenecks with this choice.

##### 5.1.1.1 VMAgent

The VMAgent component pulls metrics from exporters installed on remote hosts. The list of targets is defined by its own configuration. Scraped metrics are then sent to the VMSingle component.

#### 5.1.1.2 VMSingle

The VMSingle component is the central time series database of VictoriaMetrics. This stores collected metrics and exposes them via an HTTP API. This single service is basically the cluster version's 3 components (`vminsert`, `vmstorage`, `vmselect`) compiled into one.

#### 5.1.1.3 VMAlert

The VMAlert component reads metrics from the VMSingle component. Based on its configured rules it decides whether the conditions to fire an alert are met. If an alert fires it is sent to the VMAlertmanager component.

If recording rules are configured it is able to write the results of complex queries back to VMSingle. This is useful if we want to precalculate time consuming queries which can later be accessed just like regular metrics.

#### 5.1.1.4 VMAlertmanager

The VMAlertmanager component receives alerts from the VMAlert component. Then based on its alert routing and receivers configurations it decides if a notification should be sent. If conditions are met a notification is sent out to one or more external services, most importantly to CESIG's Mattermost (not pictured) server.

### 5.1.2 Promtail

Promtail receives and processes logs from remote hosts. Unlike VictoriaMetrics remote hosts push their logs to Promtail. Promtail then sends the processed logs to Loki.

### 5.1.3 Loki

Loki indexes and stores logs received from the Promtail component. It receives requests from Grafana, queries its internal datastore and responds with the results. Just like the VictoriaMetrics TSDB Loki can be installed as multiple instances of different components or as one single component. The former is useful for higher volumes of logs while the latter is much more simple to set up and maintain at the cost of scalability. We have opted for the simpler method as we do not expect to hit any bottlenecks here either.

### 5.1.4 Grafana

Grafana is only used for visualizations in our architecture. This is our main user-facing component: users get authenticated and authorized here. After logging in they are able to create, view and modify dashboards on Grafana. Grafana queries Loki and/or VMSingle for the data required by the currently viewed dashboard(s).

## 5.2 Deployment

All of the above listed components can be installed in a wide variety of ways. If needed these can be downloaded separately, compiled from source and then run as classic Unix

daemons. Docker images are also available for each of these. On the other end of the spectrum these components can easily run on Kubernetes as vast amounts examples, guides and official configurations are available for this kind of setup.

### 5.2.1 Kubernetes

Kubernetes[7] is an open-source system for automating the deployment, scaling, and management of containerized applications. It is a popular platform for running microservices and other distributed applications, and is often used in cloud-native environments. Kubernetes provides a number of features to help you manage your applications, including automatic scheduling and scaling of containers, self-healing, and service discovery.

We have decided to opt for Kubernetes because:

- freely available documentations and configurations make it very easy to use
- Kubernetes handles the lifecycle of jobs and pods for us
- as CESIG 's members already operate a k8s cluster they have experience using this system
- by using Kubernetes we can easily expand the system with new physical servers in the future

All in all by installing our software stack on Kubernetes we are able to implement this architecture much quicker, easier and without the need for custom tools.



## Chapter 6

# Implementation details and experiences

With the requirements and the system architecture defined we have a clear path ahead ourselves. There are still quite a few decisions to be made and problems to be solved. We are going to explore these in this chapter.

### 6.1 The server hardware and the host system

#### 6.1.1 The Hardware

The hardware was inherited from previous monitoring projects. The machine is a ProLiant DL360e Gen8 with two Intel Xeon E5-2450L CPUs, 32GBs of RAM, four Gigabit Ethernet ports, 2x120GB SSDs and 2x1TB HDDs. This should be perfect to serve our needs.

#### 6.1.2 The Operating System

For the host operating system we are going to use a Linux distribution aimed at servers. We have several choices here favouring ones that CESIG already operate. These are: Ubuntu, Debian and Rocky Linux.

Ubuntu is developed and maintained by Canonical, a UK-based company. It had been a popular choice among desktop and server Linux users alike. Unfortunately it's most recent developments such as snaps made many previous users dislike this system and migrate to others.

Debian is developed and maintained by a community of volunteers from around the world. It is known for its stability and reliability, making it a popular choice for servers and other mission-critical systems. It also has a large repository of pre-compiled software packages, making it easy to install and manage applications. Ubuntu is also based on Debian.

Rocky Linux is based on the codebase of Red Hat Enterprise Linux (RHEL), but is developed and maintained by a community of volunteers. It aims to provide a secure and stable platform that is compatible with RHEL, while also being free of any vendor restrictions or licensing fees. It also comes with SELinux preconfigured, making it more secure.

### 6.1.3 The Kubernetes Distribution

As for a Kubernetes distribution we of course have k8s. K8s supports everything we need, however it is aimed at larger installations and multi-node clusters thus it is harder to configure and more resource-intensive. A better alternative is k3s, which is a lightweight, open-source Kubernetes distribution that is designed to be easy to install and operate on resource-constrained environments, such as edge computing or IoT devices. K3s also supports SELinux and works perfectly with Rocky Linux.

### 6.1.4 The Final Host System Configuration

For the above reasons we decided to use Rocky Linux as the host OS, k3s as the Kubernetes distribution and we are going to install VictoriaMetrics and Grafana on top of these.

We installed the host system on the two SSDs configured as Linux Software RAID1 for redundancy and better uptime. The data is going to be stored on the two HDDs for which we used ZFS and the HDDs are configured as a two-way mirror (equivalent to a RAID1 in redundancy).

ZFS supports many advanced features besides different redundancy levels such as encryption, checksumming, compression and snapshots. Although we do not yet make use of these we have the opportunity to enable them in the future without worrying about migration of the system which would cause unnecessary downtime.

Another issue we need to take care of is reproducibility. As setting up such a system takes much time, effort and research it's not enough to only document this progress. During migration to new hardware (in case of an upgrade or a catastrophic failure) we would have to repeat all this work. To solve this we chose to use Ansible which is already widely used by CESIG .

Ansible is an open-source IT automation tool that allows users to configure and manage systems, deploy applications, and automate processes. It uses YAML files just like the rest of our tools, and then uses a series of modules to enforce that state on the system.

Every change is going to be applied through Ansible and every file used is going to be stored on a repository on CESIG 's GitLab server.

## 6.2 Preparing the k3s Kubernetes distribution

On it's own, Kubernetes is a very powerful tool. It can handle many containers, services, load balancers and much more. But even a single pod contains at least one docker container. It's configuration is managed via YAML files but this way of management becomes repetitive very quickly (imagine repeating the same port numbers over many configuration files) violating the DRY (Don't Repeat Yourself) principle and becomes a nightmare to maintain. To overcome this issue the community had created several tools to ease it's use.

### 6.2.1 Management Tools

The k3s kubernetes distribution can be managed the same way as it's heavyweight, full-fledged k8s brother. Thus we are going to use the same tools we'd use on a more complicated deployment.

Managing a kubernetes cluster can be (and I usually done) remotely. By this I don't mean an SSH session - the Kubernetes command-line tool, `kubectl` tool runs on any administrator's machine and for authentication we need to copy the `/etc/rancher/k3s/k3s.yaml` configuration file from the k3s host machine.

Now we have the power to manage our cluster yet this does not free us from writing complicated, thousands-of-lines long YAML configuration files for deployment. The solution is Helm.

Helm is an open-source package manager for Kubernetes that allows users to easily install, upgrade, and manage applications on Kubernetes clusters. Helm uses a packaging format called charts, which are pre-configured templates that define the desired state of an application, including its dependencies, configuration, and resources.

### 6.2.2 Persistent Volumes

VictoriaMetrics' and Loki's stateful components are going to need a persistent volume which we want to store on our ZFS mirror. Persistent volumes are going to be provided by the OpenEBS ZFS CSI Driver. Installation is easy: `kubectl apply -f https://openebs.github.io/charts/zfs-operator.yaml`.

We have to configure the driver to use our pool.

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: openebs-zfspv
parameters:
  recordsize: "4k"
  compression: "off"
  dedup: "off"
  fstype: "zfs"
  poolname: "spinningrust/openebs-zfspv"
provisioner: zfs.csi.openebs.io
```

**Listing 6.1:** zfs-localpv.sc.yaml

Configuration is applied yet again by the following command: `kubectl apply -f zfs-localpv.sc.yaml`.

## 6.3 Installing and configuring a basic VictoriaMetrics on k3s

Now that we have set up k3s we can begin installing and configuring VictoriaMetrics.

Before we continue let's take a look at how Helm works. Complex, multi-pod kubernetes systems are set up by many large YAML files. These files contain redundant informations, hard to comprehend and difficult to work with. Helm uses template files (which use Go syntax) which handle everything that can be computed, and so called values files. The latter contain sane default configurations where such things make sense and require the user to enter values where no such configurations can exist, for example external services' authentication details. The combination of these is called a helm chart. Helm charts can be added to the system manually or by adding a helm chart repository.

### 6.3.1 Downloading the Helm Charts

First of all we are going to add the VM helm chart repository following the instructions on VM's GitHub repository: <https://github.com/VictoriaMetrics/helm-charts>.

```
helm repo add vm https://victoriametrics.github.io/helm-charts/  
helm repo update
```

Next, we are going to select the helm chart that fits our needs. The repository contains many charts most of which only contain a subset of a full VictoriaMetrics stack. The one we need is `victoria-metrics-k8s-stack`, this is the full k8s VM stack.

Before installing this we have to configure it through the correct values file. We can get the values file with the help of the `helm show values` command:

```
helm show values vm/victoria-metrics-k8s-stack > values.yml
```

For the initial `values.yml` see Appendix A.1.

### 6.3.2 Installing VictoriaMetrics

Once the `values.yml` is configured the helm chart is installed with the following command:

```
helm install beholder vm/victoria-metrics-k8s-stack -f values.yml
```

## 6.4 Collecting metrics

With VictoriaMetrics up and running we can begin collecting our metrics from various servers and network devices. In fact, we are already collecting some metrics: the helm chart used before already configured the monitoring of VictoriaMetrics, k3s and the host node itself.

The stack installed also contains another powerful feature: the VictoriaMetrics Kubernetes Operator.

An operator is a piece of software that extends the Kubernetes API to create, configure, and manage instances of complex applications. This allows you to use the Kubernetes API to automate tasks such as deploying, scaling, and backing up your applications. They provide a declarative way to describe the desired state of an application, and then automatically reconciling the actual state of the application with the desired state.

### 6.4.1 Collecting Metrics from Servers

We are going to add 3 servers for demonstration purposes:

- memory-alpha and memory-beta: this is a ZFS-based storage server cluster
- trash: this server provides music for CESIG 's hard-working members

To collect metrics from a remote system the target host must have the necessary exporters installed and configured. Exporters are a piece of software that collect metrics from system components and expose them through an HTTP API. Different kinds of exporters collect different metrics: the generic node exporter collect basic statistics such as CPU and RAM usage, while specialized exporters are needed to collect metrics from components such as ZFS, mailing servers or databases.

After the desired exporters are installed and configured on the target host we have to inform VictoriaMetrics about their existence. For this we have to define `VMStaticScrape` objects, apply these configurations and let the operator configure the `vmagent` component to collect metrics from these servers as well. The configuration is as follows:

```
apiVersion: operator.victoriametrics.com/v1beta1
kind: VMStaticScrape
metadata:
  name: memory-alpha-and-beta
  namespace: monitoring
spec:
  targetEndpoints:
    - targets:
      - 10.0.69.101:9100
      - 10.0.69.102:9100
  jobName: memory-alpha-and-beta
```

**Listing 6.2:** Configuration file for memory-alpha and memory-beta

```
apiVersion: operator.victoriametrics.com/v1beta1
kind: VMStaticScrape
metadata:
  name: trash
  namespace: monitoring
spec:
  targetEndpoints:
    - targets:
      - 152.66.209.79:9100
  jobName: trash
```

**Listing 6.3:** Configuration file for trash

The `.metadata.name` and `.metadata.namespace` values identify the resource, while key-value pairs under the `.spec` key define the configuration itself. The `.spec.targetEndpoints.targets` array contain the ip+port pairs for every available exporter, while `.spec.jobName` is used to locate the collected metrics inside `vmagent` and `Grafana`.

These files contain many boilerplate around little information. Luckily we have a solution for this which we going to explore in the next section.

The monitoring server has a dedicated interface in the administration VLAN just like every other server operated by CESIG - this means that we can reach these targets directly.

## 6.4.2 Collecting Metrics from Network Devices

Monitoring network devices is slightly more complicated. Network devices can not run arbitrary softwares such as exporters and these devices can only be reached on a physically separated network knowns as the management LAN. Luckily we can reach the NOCs (Network Operation Centers) which are connected to this network. For the other part we must use some kind of exporter that runs separately and collects data through SNMP.

The Prometheus project has a solution exactly for this: [https://github.com/prometheus/snmp\\_exporter/tree/main/generator](https://github.com/prometheus/snmp_exporter/tree/main/generator) and CESIG already has an SNMP exporter with the required MIBs built in. We installed one instance of this exporter on

each NOC. The exporters are reached through a k3s loadbalancer, this makes sure that the metrics are not duplicated and so far at least one exporter is running metrics are going to be collected.

```
apiVersion: v1
kind: Service
metadata:
  name: noc-snmp-exporter
spec:
  ports:
    - port: 80
      targetPort: 9116
      name: exporter
      protocol: TCP
  type: ClusterIP
```

---

```
apiVersion: v1
kind: Endpoints
metadata:
  name: noc-snmp-exporter
subsets:
  - ports:
      - port: 9116
        name: exporter
        protocol: TCP
    addresses:
      - ip: 10.0.208.241
      - ip: 10.0.208.242
```

**Listing 6.4:** Load balancer configuration

With all these set up we can scrape network devices just like we did with servers. Below is showed the initial configuration scraping the core router `rtr-1` only.

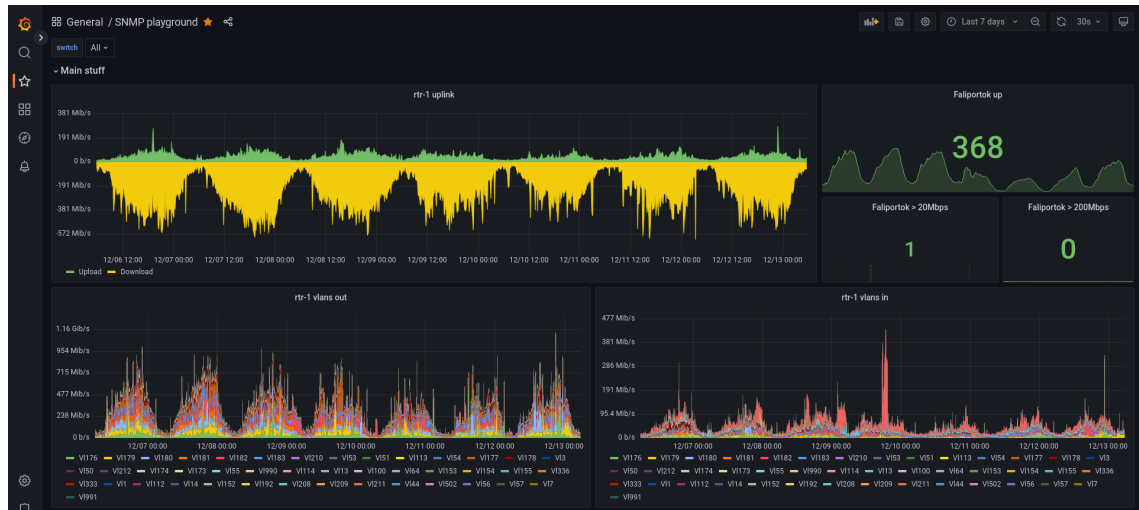
```
apiVersion: operator.victoriametrics.com/v1beta1
kind: VMStaticScrape
metadata:
  name: schnet-snmp
  namespace: monitoring
spec:
  jobName: schnet-snmp
  targetEndpoints:
    - targets:
        - noc-snmp-exporter:80
      path: /snmp
      params:
        target: [ rtr-1 ]
        module: [ sch_core_net_mibs ]
```

**Listing 6.5:** Initial SNMP scraping configuration

Note that all network device metrics are reached through a single exporter. The `.spec.targetEndpoints.params.target` key identifies the device.

### 6.4.3 Metric Visualizations

As for visualization, our **Grafana** comes with some common dashboards preinstalled such as a node exporter, kubernetes exporter and VictoriaMetrics dashboards. Other, commonly used dashboards can easily be imported from <https://grafana.com/grafana/dashboards/>. Unique dashboards, such as the ones displaying our SNMP metrics must be created by us.



**Figure 6.1:** Screenshot of the Grafana SNMP Playground Dashboard at the time of writing

The above is shown a screenshot of the SNMP Playground dashboard. After the initial version was released members of CESIG took their time and added many other useful metrics to this visualization. The 'Faliportok up' panel shows every active 1Gbit endpoint on the dormitory's network. It shows nicely the least actively used periods of the week: between 4am and 8am especially on the weekends.

## 6.5 Automating tasks with the help of GitLab CI

Now we are able to monitor our servers and network devices. However as seen above, configuration is very repetitive. Moreover, adding or modify any device requires the contribution of the monitoring system administrator which slows down the integration and wastes time that could be utilized better elsewhere.

We already have a perfect tool for generating configurations: helm charts. We can quickly and easily create simple ones ourselves. For defining a helm chart we need:

- `./Chart.yaml`: this contains some required metadata about the chart itself
- `./template/*.yaml`: these are the Go templates themselves, the final config will be generated from these files
- `./template/*.tpl`: these are helper template functions
- `./values.yaml`: the values for the template

### 6.5.1 Implementing Helm Charts for Configuration Generation

Let's review the implementation of the node exporter helm chart.

```
apiVersion: v2
name: node-exporter-target
description: Scrape configs for Prometheus node exporter
type: application
version: 0.0.0
appVersion: "kolbaasz"
```

Listing 6.6: ./Chart.yaml

```
{{- define "name" }}
{{- .Chart.Name }}-{{.Release.Name}}
{{- end }}
```

Listing 6.7: ./templates/\_helpers.tpl

```
{{- $global := . }}
{{- range .Values.Devices }}

apiVersion: operator.victoriametrics.com/v1beta1
kind: VMStaticScrape
metadata:
  name: "{{ template "name" $global }}-{{ .Name }}"
  namespace: monitoring
spec:
  targetEndpoints:
    - targets:
        {{- range .Targets }}
        - "{{ . }}"
        {{- end }}
    jobName: "{{ .Name }}"
{{- end }}
```

Listing 6.8: ./templates/node-exporter.yaml

```
Devices:
- Name: memory-alpha-and-beta
  Targets:
    - 10.0.69.101:9100
    - 10.0.69.102:9100

- Name: trash
  Targets: [ 152.66.209.79:9100 ]
```

Listing 6.9: ./values.yaml

These Go templates operate with the values file in context as `.Values`. This chart iterates over the `.Values.Devices` array's elements with the `range` function, creates a `VMStaticScrape` object for each one. Each `VMStaticScrape`'s `.spec.targetEndpoints.targets` array is filled with the matching `.Values.Devices.Targets` array's elements.



Now when adding a new device the administrator only has to edit the `values.yml` file which now only contains the necessary informations without any noise.

To start using this new method, first we have to install the helm chart: `helm install -n monitoring node-exporter-targets ..` For subsequent uses we can upgrade the chart: `helm upgrade -n monitoring node-exporter-targets ..`

We can generate SNMP monitoring configurations the same way. See the helm chart in Appendix A.2.

### 6.5.2 Configuring GitLab CI

We had achieved a way to write monitoring configurations easily but this approach still requires the intervention of a precious sysadmin. Continuous integration is our way to go. Our repository on CESIG's GitLab contains everything needed to add new hosts to the system, let's make use of GitLab's Continuous Integration (CI) features. For this the following `.gitlab-ci.yml` file is going to be used.

```
stages: [ sanity_check, upgrade_charts ]

before_script:
  - apk add diffutils
  - kubectl config use-context kszk/ansible/playbooks/monitoring:beholder

node_check:
  stage: sanity_check
  image: alpine/k8s:1.24.8
  script:
    - ./helmdiff.sh node-exporter-targets ./files/helm-charts/node-exporter-target
  only:
    changes:
      - files/helm-charts/node-exporter-target/**
      - .gitlab-ci.yml
      - helmdiff.sh

snmp_check:
  stage: sanity_check
  image: alpine/k8s:1.24.8
  script:
    - ./helmdiff.sh snmp ./files/helm-charts/snmp-exporter-target
  only:
    changes:
      - files/helm-charts/snmp-exporter-target/**
      - .gitlab-ci.yml
      - helmdiff.sh

node_export:
  stage: upgrade_charts
  image: alpine/k8s:1.24.8
  script:
```

```

- helm upgrade -n monitoring node-exporter-targets ./files/helm-
charts/node-exporter-target
only:
  changes:
    - files /helm-charts/node-exporter-target/**
    - .gitlab-ci.yml
    - helmdiff.sh
  refs:
    - main

snmp_export:
  stage: upgrade_charts
  image: alpine/k8s:1.24.8
  script:
    - helm upgrade -n monitoring snmp ./files/helm-charts/snmp-
exporter-target
  only:
    changes:
      - files /helm-charts/snmp-exporter-target/**
      - .gitlab-ci.yml
      - helmdiff.sh
    refs:
      - main

```

**Listing 6.10:** .gitlab-ci.yml

Note the `helmdiff.sh` script used above. We tried to use the already available helm diff plugin but ran into several problems, the version we used did not accept our new and completely valid configurations. For this reason we used the following script:

```

#!/bin/sh

set -ueo pipefail

release="$1"
chart="$2"

helm get manifest -n monitoring "$release" > /tmp/installed
helm template -n monitoring "$release" "$chart" > /tmp/new

diff -w --color=always -y /tmp/installed /tmp/new || [ "$?" = "1" ]

```

**Listing 6.11:** helmdiff.sh

The CI scripts run on the monitoring server in docker inside an Alpine image. The image is complemented with the `diffutils` package which is required by the `helmdiff.sh` script. The CI runs on two stages for each of our helm charts. The first stage, `sanity_check` verifies the syntactic validity of our new configuration as well as displays the difference between the current and the newly generated difference using `helmdiff.sh`. On the main branch the `upgrade_charts` stage upgrades our helm charts using the new configurations. One job per helm chart is used in each stage and each job is only ran when necessary (i.e. when a change was made).

Using these tools new systems can be monitored with minimal knowledge and without the intervention of a monitoring system administrator.

## 6.6 Authentication and authorization

So far our users can set up monitoring for any of their devices but we are far from done. As only the hardcoded **Administrator** account can log into the otherwise publicly available Grafana users can not make use of the system.

We are going to authenticate and authorize our users through AuthSCH with OAuth 2.0. Luckily Grafana supports this protocol. Unlucky for us Grafana expects the user's e-mail to be at the `/email` path exactly and we found no way to reconfigure this. For legacy reasons AuthSCH serves the e-mail in a different path.

To work around this issue we created a very simple proxy server that translates AuthSCH's response to the format expected by Grafana. Since the implementation of this service is trivial we are not going to explore this further.

```
grafana:
  grafana.ini:
    auth.generic_oauth: # client is at Garmine's SCHAcc
      name: "AuthSCH"
      enabled: "true"
      client_id: "18023005984133885935"
    client_secret: "secret" # Redacted for security purposes
    scopes: "basic displayName admembership linkedAccounts"
    auth_url: "https://auth.sch.bme.hu/site/login"
    token_url: "https://auth.sch.bme.hu/oauth2/token"
    api_url: "http://authsch-api-authsch.authsch-hack.svc/grafana" # ha-ha
    allow_sign_up: "true"
    role_attribute_path: "contains(adgroups[*], 'KSZK') && 'Admin' ||
contains(adgroups[*], 'KSZKUjonc') && 'Viewer' || 'Vesztettem'"
    role_attribute_strict: "true"
    groups_attribute_path: "adgroups"
```

Listing 6.12: Grafana OAuth 2.0 configuration

The above is a pretty standard OAuth 2.0 configuration. Notable are the different `api_url` path (these requests are routed through our proxy) and the `role_attribute_path` value. The latter describes the authorization based on AD group memberships. Members of the **KSZK** group are granted **Admin** role: they are authorized to do make any changes in Grafana save for a few things such as managing users and changing permissions. Members of the **KSZKUjonc** group are granted **Viewer** role: they are authorized to view almost any part of our Grafana but can not make any changes. Other users would get **Vesztettem** which is a nonexistent role, thus they are denied logging in.

Superadmin rights are granted locally either through the above mentioned **Administrator** account or by another superadmin account.

## 6.7 Collecting logs

### 6.7.1 Configuring and Installing Loki

We are going to use Loki for writing, storing and querying logs. Loki just like Victoria-Metrics can be split into smaller services and installed in a larger, distributed and scalable form as a cluster for heavier workloads. For smaller infrastructures Loki can be installed

as a single service as well which is easier to configure and should fit our needs perfectly. For this we are going to write Loki's `loki-values.yml` file.

First of all we are going to install the helm chart as before:

```
helm repo add grafana https://grafana.github.io/helm-charts
helm repo update
```

**Listing 6.13:** Installing Loki's helm chart

```
loki:
  commonConfig:
    replication_factor: 1
  storage:
    type: filesystem
  server:
    log_level: warn

monitoring:
  selfMonitoring:
    enabled: false

test:
  enabled: false

write:
  replicas: 0

read:
  replicas: 0

singleBinary:
  replicas: 1
  persistence:
    size: 256Gi
    storageClass: openebs-zfspv

ingress:
  enabled: false

gateway:
  enabled: false

networkPolicy:
  enabled: false

minio:
  enabled: false
```

**Listing 6.14:** `loki-values.yml`

In order to configure Loki as a single binary we need a single replica of `singleBinary` and no `write` and `read` replicas. Logs are going to be stored on a persistent volume provided by `openebs-zfspv`. For now we are going to disable `selfmonitoring` and `testing` as well.

Once the configuration is ready we are going to install as usual: `helm install --values loki-values.yaml loki grafana/loki`

Loki only receives logs from another service and does no processing on them. In order to send logs to Loki we are going to use Promtail. Promtail is very versatile, is able to receive and collect logs from various sources and can process the logs before sending them. We are going to split the responsibility of collecting and processing logs. Logs are going to be collected and sent by the hosts themselves either via Promtail (for servers) or syslog (for more simple devices such as routers and switches). Processing the logs is going to be done by a Promtail instance running on the monitoring system. This way target hosts can send logs with minimal and very simple configurations while the monitoring system administrators can configure log processing locally, without having access to the target hosts.

### 6.7.2 Collecting Logs of the Monitoring System

The first target for log collecting is the monitoring system itself. So far only the system administrators had access to these logs via the `kubectl logs` command. Having access to these logs would make debugging much easier for our users.

```
daemonset:
  enabled: false

deployment:
  enabled: true
  replicaCount: 1

serviceMonitor:
  enabled: true

config:
  logLevel: warn
  serverPort: 3101
  clients:
    - url: http://loki:3100/loki/api/v1/push
  snippets:
    common:
      - action: replace
        target_label: job
        replacement: beholder-pods
```

**Listing 6.15:** local-promtail-values.yml

The local Promtail's configuration is changed minimally: since we are running a single node kubernetes cluster we opt for a single replica deployment, `serviceMonitor` is enabled for the VictoriaMetrics operator and collected logs' are stored under the 'beholder-pods' job for easy identification.

### 6.7.3 Receiving and Processing Logs from other Systems

The other Promtail instance is going to receive, process and then send the remotely collected logs to Loki. Let's take a look at its configuration.

```

daemonset:
  enabled: false

deployment:
  enabled: true
  replicaCount: 1

defaultVolumes: []
defaultVolumeMounts: []

serviceMonitor:
  enabled: true

extraPorts:
  syslog:
    name: tcp-syslog
    containerPort: 1514
    protocol: TCP
    externalTrafficPolicy: Local
    service:
      type: ClusterIP # TODO make LoadBalancer after auth has been figured
      out
      port: 1514
  promhttp:
    name: prom-http
    containerPort: 3500
    protocol: TCP
    service:
      type: ClusterIP
      port: 3500
  promgrpc:
    name: prom-grpc
    containerPort: 3600
    protocol: TCP
    service:
      type: ClusterIP
      port: 3600

config:
  file: |
    server:
      log_level: info
      http_listen_port: 3101
    clients:
      - url: http://loki:3100/loki/api/v1/push

    positions:
      filename: /run/promtail/positions.yaml

  scrape_configs:

```

```

- job_name: syslog
  syslog:
    listen_address: 0.0.0.0:1514
    labels:
      job: syslog
  relabel_configs:
    - source_labels: [ __syslog_message_hostname ]
      target_label: host
    - source_labels: [ __syslog_message_severity ]
      target_label: severity
    - source_labels: [ __syslog_message_facility ]
      target_label: facility
    - source_labels: [ __syslog_message_app_name ]
      target_label: app

- job_name: prom
  loki_push_api:
    server:
      http_listen_port: 3500
      grpc_listen_port: 3600
    labels:
      job: prom

limits_config: {}

```

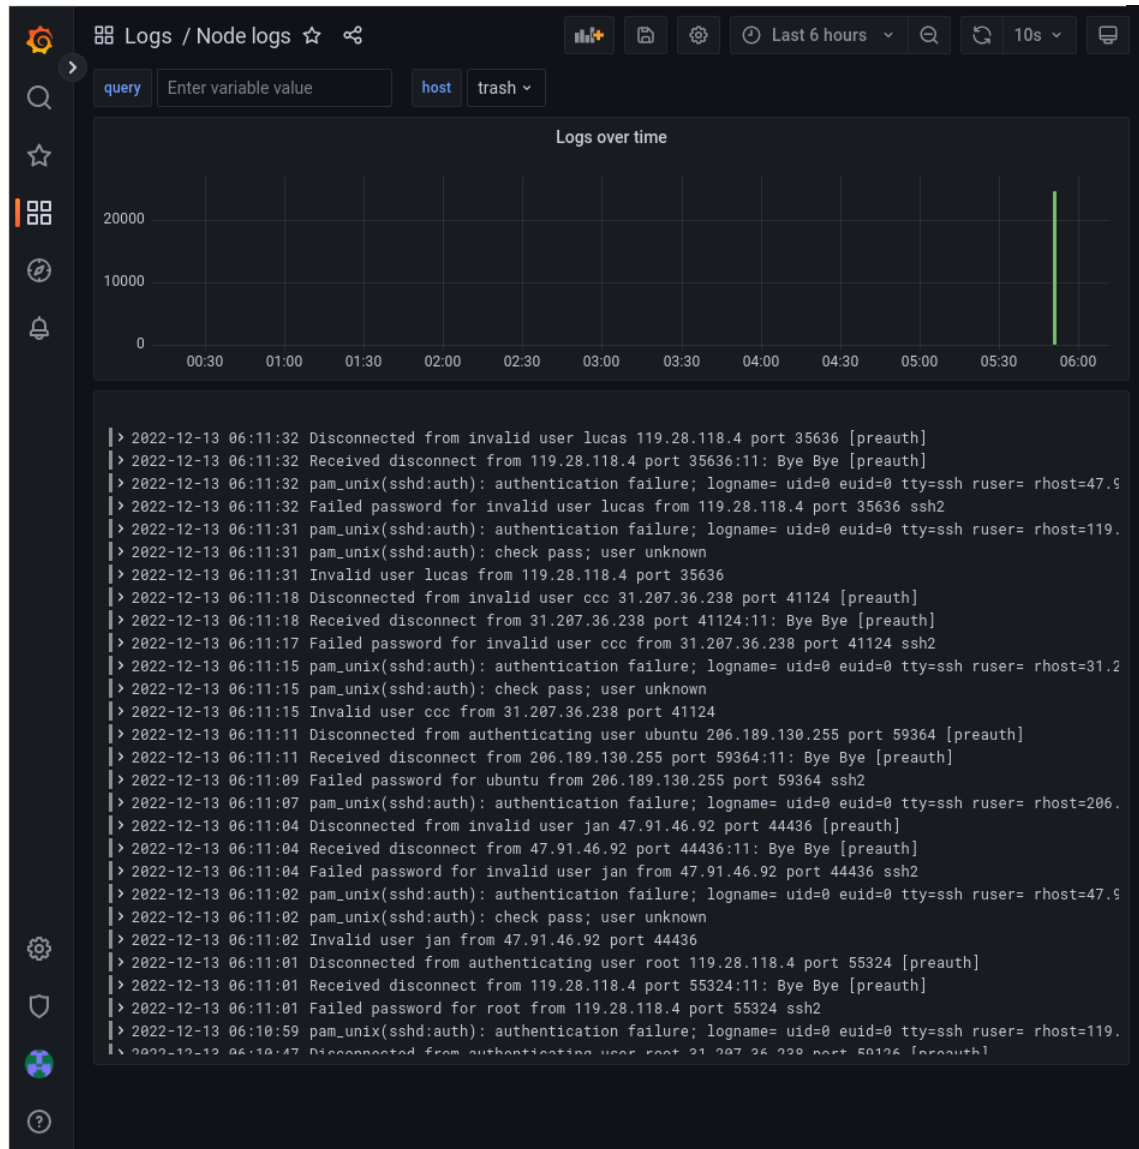
**Listing 6.16:** promtail-values.yml

Similarly to the local Promtail we opt for a single replica deployment. As we do not wish to collect local logs with this instance we disable `defaultVolumes` and `defaultVolumeMounts`. The `extraPorts` map configures Promtail to accept connections from syslog and other Promtail instances. The `config.file` is just a string, literally the contents of Promtail's configuration file. The above is just a simple forwarding for Promtail and syslog, with added label translation for syslog.

Installation is done with `helm install` for both charts just as before.

For demonstration purposes we installed a Promtail instance on a server called `trash`. After setting up a Grafana dashboard we are able to successfully browse `trash`'s logs.

## 6.7.4 Log Visualization



**Figure 6.2:** Screenshot of the Grafana Logs dashboard just after implementation

Above is seen a screenshot of the Grafana Logs dashboard. On the logs collected from Trash, CESIG's media-player VM we can see several unsuccessful log-in attempts from various botnets.

## 6.7.5 Security

For security purposes authentication is mandatory for receiving logs. Both Promtail->Promtail and syslog->Promtail ways are able to use authentication. As many network devices don't support anything else than sending logs via unauthenticated and unencrypted syslog we have postponed this feature.



## 6.8 Creating alerts and sending notifications

In the monitoring stack the `vmalert` component communicates with the VictoriaMetrics TSDB, has access to every metric and executes recording and alerting rules. Recording rules' results are written back to the TSDB while alerting rules send notifications to the `vmalertmanager` component. The alert manager then routes these alerts to various receivers. Receivers then finally format and send out alerts to external platforms. These include but are not limited to: e-mail, Slack API, Telegram and HTTP.

### 6.8.1 Configuring VMAAlert

Our `vmalert` configuration is very plain. The bulk of it's configuration, the alerting and recording rules are applied through the VictoriaMetrics kubernetes operator.

```
vmalert:
  annotations: {}
  enabled: true
  spec:
    selectAllByDefault: true
    image:
      tag: v1.83.0
    evaluationInterval: 15s

  templateFiles:
    {}

  ingress:
    enabled: false
```

**Listing 6.17:** vmalert configuration

Luckily the kubernetes VictoriaMetrics stack comes with a plethora of alerting rules pre-configured. Some of these alerts only fire on VictoriaMetrics' error conditions while others are applied for every monitored node such as NTP out of sync and RAID error alerts.

```
apiVersion: operator.victoriametrics.com/v1beta1
kind: VMRule
metadata:
  name: beholder-victoria-metrics-k8s-stack-node-exporter
  namespace: monitoring
spec:
  groups:
    - name: node-exporter
      rules:
        - alert: NodeClockSkewDetected
          annotations:
            description: Clock is out of sync by more than 300s. Ensure NTP is configured correctly on this host.
            runbook_url: https://runbooks.prometheus-operator.dev/runbooks/node/nodeclockskewdetected
            summary: Clock skew detected.
          expr: |-
```

```

    ( node_timex_offset_seconds > 0.05 and deriv(
node_timex_offset_seconds[5m]) >= 0 )
    or
    ( node_timex_offset_seconds < -0.05 and deriv(
node_timex_offset_seconds[5m]) <= 0 )
  for: 10m
  labels:
    severity: warning

```

**Listing 6.18:** Example alerting rule configuration

The above is an example for alerting when the system clock drifts more than 300 seconds on any monitored node.

## 6.8.2 Configuring VMAalertmanager

These alerting rules are ideal for demonstration purposes and are also useful for real-life scenarios. But with nowhere to go they would be worthless, so let's configure our alert manager to send everything to CESIG's Mattermost server's `~alerts-test` channel for now.

```

alertmanager:
  enabled: true
  spec:
    selectAllByDefault: true
    externalURL: ""
    routePrefix: /

  config:
    global:
      resolve_timeout: 5m
      slack_api_url: "https://127.0.0.1" # Redacted for security reasons

    templates:
      - "/etc/vm/configs/**/*.tmpl"

    route:
      group_by: ["alertgroup", "job"]
      group_wait: 30s
      group_interval: 5m
      repeat_interval: 12h
      receiver: "mattermost"

  receivers:
    - name: "mattermost"
      slack_configs:
        - channel: "#alerts-test"
          username: "SCP-049"
          send_resolved: true
          title: '{{ template "slack.monzo.title" . }}'
          icon_emoji: '{{ template "slack.monzo.icon_emoji" . }}'
          color: '{{ template "slack.monzo.color" . }}'
          text: '{{ template "slack.monzo.text" . }}'

```

```

    actions:
      - type: button
        text: "Runbook :green_book:"
        url: "{{ (index .Alerts 0).Annotations.runbook_url }}"
      - type: button
        text: "Dashboard :grafana:"
        url: "{{ (index .Alerts 0).Annotations.dashboard }}"
      - type: button
        text: "Silence :no_bell:"
        url: '{{ template "__alert_silence_link" . }}'

# better alert templates for slack
# source https://gist.github.com/milesbxf/
# e2744fc90e9c41b47aa47925f8ff6512
monzoTemplate:
  enabled: true

templateFiles:
  {}

ingress:
  enabled: false

```

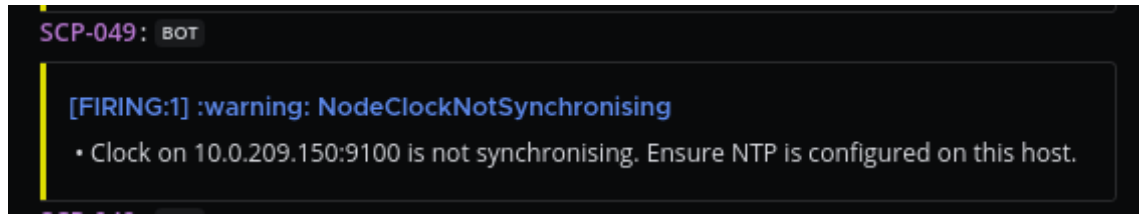
**Listing 6.19:** alertmanager configuration

Mattermost uses the Slack API for incoming webhooks. This sample `alertmanager` config uses trivial routing. A single failure usually triggers more alerts at once so it only tries to send alerts as a batch in a single message instead of one-by-one in order to reduce noise. The receiver config on the other hand highlights some nice features:

- we receive a notification when the cause is fixed
- the monzo templates nicely decorate the message
- we get a button to the relevant runbook (basically a manual for the alert)
- we get a button to the relevant Grafana Dashboard
- we get a button that silences the alert

### 6.8.3 Alerts on Mattermost

After applying the new configuration using `helm upgrade` we instantly receive several alerts on Mattermost's `~alerts-test` channel that had been stuck in the queue.



**Figure 6.3:** Screenshot of a real alert on CESIG 's Mattermost server

Above is shown a real alert sent from VictoriaMetrics to CESIG 's Mattermost server. This alert unexpectedly unveiled a real problem: one of our servers system clock did not properly synchronise via NTP. The difference was around 500 milliseconds.

As it turns out the Mattermost API is not 100% compatible with Slack's: even though both support interactive messages and buttons we do not get them with the alert messages.

### 6.8.4 Implementing Continuous Integration for Alerts

Even though the preconfigured alerts are useful we want to give out users the ability to define their own alerts. For this we are going to use GitLab CI as well.

Ideally we would have a single helm chart for everything: monitoring, logging and alerting configurations. This would eliminate every repetition and with a preconfigured rules library would make alerting much more simple. This would require lots of work and careful planning. This is out of the scope for this project and thus we have opted for a much more trivial solution. See the helm chart for adding alerting and recording rules in Appendix A.2.

We installed the helm chart usnig the `helm install` command and added the following new jobs to `.gitlab-ci.yml`

```
alerts__check:
  stage: sanity__check
  image: alpine/k8s:1.24.8
  script:
    - ./helmdiff.sh kszk-alerts ./files/helm-charts/kszk-alerts
  only:
    changes:
      - files/helm-charts/kszk-alerts/**
      - .gitlab-ci.yml
      - helmdiff.sh

alerts__upgrade:
  stage: upgrade_charts
  image: alpine/k8s:1.24.8
  script:
    - helm upgrade -n monitoring kszk-alerts ./files/helm-charts/kszk-alerts
  only:
```

```
changes:
  - files /helm-charts/kszk-alerts/**
  - .gitlab-ci.yml
  - helmdiff.sh
refs:
  - main
```

**Listing 6.20:** New jobs in .gitlab-ci.yml

Although it's more complicated than ideal we gave our users the tools necessary to configure their own alerting and recording rules without the intervention of a monitoring system administrator.

# Chapter 7

## Results and evaluation

During the time of writing the monitoring system has been running for months now with gradually increasing capabilities as time went on. Although most requirements were met the system is still in development. Let's take a look at the current state of it.

### 7.1 Evaluation of Requirements Fulfillment

The most important is meeting the requirements laid out in Chapter 3.

#### 7.1.1 Collecting and storing metrics and logs

The currently installed system is more than capable of collecting and storing metrics and logs. We are able to monitor every server platform and network device currently in use by CESIG .

At the time of writing many systems including Windows Servers, Linux Servers and Cisco devices are already being monitored. Only logs are not being collected from network devices due to debates about an optimal and secure solution.

#### 7.1.2 Effective long-term storage of historical data

The system is currently capable of storing data for short to mid terms. The implementation of long-term storage was least concern. With proper configuration such as fine tuned retentions and external storage both VictoriaMetrics and Loki are able to effectively store large amounts of data only limited by available storage.

#### 7.1.3 Detecting issues and alerting the necessary personnel

VictoriaMetrics' components already detect some basic issues and send out alerts to a test channel on CESIG 's Mattermost platform. The software stack is capable of detecting much more complex issues and sending out alerts on many communications channels. This perfectly demonstrates our top priority: preventing future issues and detecting existing ones.

#### **7.1.4 Detecting issues of the monitoring system itself (meta-monitoring)**

Meta-monitoring is partially implemented. The system collects and visualizes many metrics and logs about itself as well as many alerts configured to notify its maintainers about intervention needed. There are still some places to improve though. External monitoring of the system is very much possible but is not yet implemented at all. This is a crucial next step for production.

#### **7.1.5 Ease of use yet highly customizable**

The addition of CI and configuration generation made maintenance and the connecting of new systems very easy. The usage of industry-standard tools such as Kubernetes and Helm Charts and simple data formats such as YAML files help make the learning curve shallow. Internal documentations as well as this paper also help users and administrators alike.

#### **7.1.6 Authentication and authorization via AuthSCH**

Authentication and authorization via AuthSCH using OAuth 2.0 protocol is fully implemented and is already being used by CESIG's members.

#### **7.1.7 Visualization of collected data**

Visualization of collected logs and metrics is done using Grafana, a leading industry-standard tool. The vast array of publicly available dashboards and the easy configuration of new ones means it is able to handle everything collected by the rest of the software stack. The configured UI already handles the most important and interesting data correctly and presents them in an easy to understand graphical way.

#### **7.1.8 Security**

Access to the system is only granted to CESIG's members, the server is equipped with up to date software, an open source version of an enterprise Linux distribution, protected by firewall and connection can only be made through TLS. Detailed security analysis was not yet done as some features still have to be improved upon.

#### **7.1.9 Summary**

All in all the new monitoring system already fully satisfies 4 and partially satisfies 3 out of the 8 main requirements. There is no known difficulty in fully satisfying every requirement.

### **7.2 Evaluation of the Performance of the System**

#### **7.2.1 CPU and RAM Usage**

Currently system load and CPU usage is around 4% on average. RAM usage is 12GBs out of 32GB totally installed

VictoriaMetrics scales very well even with much larger amounts of collected metrics.

Loki's memory usage depends on its configurations (e.g. cluster or monolithic modes) and the types of queries performed. The current setup uses small amount of memory when executing simpler queries. The amount of available RAM should be fine even with more hosts. For more complex queries with a larger number of hosts fine tuning of Loki is going to be necessary and switching to a cluster mode of Loki is advised.

### **7.2.2 Disk Usage**

At the moment only 13% of the SSDs are being used by the host system. This is not expected to change much in the future as both metrics and logs are stored on the HDDs.

Out of the 1TB available space on the ZFS pool located on the HDDs around 5GB is already being used by metrics and logs collected. This is after around one month of continuous data collection. At this rate the HDDs will be fully utilized in 17 years. Right now at most 50% of available metrics and 10% of available logs are being collected. With larger adaptation this would only last around one to two years.

Considering this during the following year data retention policies will have to be fine tuned and new, larger capacity HDDs will have to be acquired. A long-term and more effective storage solution should be developed in the next two years.

### **7.2.3 Network Usage**

Network traffic is very minimal. The server has 4x1Gbit ethernet ports yet we only use around 4Mbit/s on the single connected port. Even with all ports redundantly configured and 500 times more incoming data the network bandwidth will not be a bottleneck.

### **7.2.4 Summary**

The currently configured storage is ideal for development and testing purposes but it's going to be the first issue mid-term. This is very much expected from a service that collects data in every millisecond in its uptime. For CESIG's infrastructure size CPU and RAM may not be a problem but fine tunings and configuration changes will have to be made on the long run.

## **7.3 Results of Continuous Integration**

### **7.3.1 CI for Collecting Metrics from Servers**

Before CI only 3 servers were monitored for demonstration purposes. After the release of CI features to CESIG's members they quickly and voluntarily started to connect their own systems with only minimal to no help other than the provided documentation.

Systems added by our users through CI:

- 11 Linux systems with 19 exporters in total
- 6 Windows systems with 6 exporters in total



Exporters installed by our users include from the generic Linux/Windows node exporter to more specific ones such as database and ZFS exporters. They also imported or created many new dashboards when it was necessary.

### **7.3.2 CI for Collecting Metrics from Network Devices**

CI was also implemented for network devices and our users made use of it. After release 11 more Cisco devices were added to the system similarly to the previous case. A basic SNMP dashboard was also released that was later extended and improved by members of CESIG .

### **7.3.3 CI for Configuring Alerts**

No new alerts were configured yet. The obvious cause of this is a lack of time as alerts were introduced recently compared to the time of writing. Alerts' configuration generation is also much less automated than the previously reviewed ones.

### **7.3.4 Summary**

The three cases discussed above prove that with a carefully designed and easy to use configuration generation system members of CESIG will voluntarily connect their systems into the new monitoring apparatus with minimal or no help from its maintainers.

## **7.4 Other Results of the Monitoring System running so far**

The self-monitoring capabilities already helped during some experimentation and debug sessions.

The preconfigured alerts managed to find two NTP synchronization issues immediately after only monitoring was configured. After the alert manager was properly configured later on we instantly received notifications for these alerts on CESIG 's Mattermost.

## Chapter 8

# Future Development Opportunities

Besides the smaller imperfections and missing features described in the previous chapter there are still many great ideas that could be implemented in the future. Some of these are detailed in this chapter.

### 8.1 Improving Configuration Management and Enforcing Reproducibility

During development it was easier to take some shortcuts. Unfortunately this has resulted in violation of reproducibility. This must be addressed in the future.

#### 8.1.1 Grafana Configurations

Grafana has a rich UI where users can configure and edit many things. This also means that any changes made are stored in some format internally. Currently these would be lost during a reinstallation. Applying some sort of version control to these would solve this problem.

#### 8.1.2 Backups

Currently there are no backups of the system. In case of a catastrophic failure data and configurations could be lost forever which is unacceptable. Regular automated backups are the solution.

#### 8.1.3 Ansible Scripts

Ansible scripts might be out of sync currently. Regularly verifying the reproducibility of installation and configuration would ensure the integrity of these. These, together with backups would make recoveries and hardware upgrades very quick and painless.

### **8.1.4 Advanced Configuration Generation**

Currently the three main configurations (exporters, SNMP and alerts) are completely separate. With the development of some common data model and configuration generation these could be unified. Adding a library of common configurations (such as alerts) would make the addition of new systems seem like child's play.

## **8.2 Granting Users Access to Internal Components**

Many information about the system's internal states are available to our users thanks to self-monitoring. Yet there are still some components that would be very useful to them for debugging and experimentation.

These are not available to users at the moment because authentication and authorization is not yet solved for them.

### **8.2.1 Access to the VMAgent component**

VMAgent has a web UI where status and informations are displayed about every exporter. In some cases VMAgent has access to a specific exporter yet it's metrics are not seen by Grafana. This would help rule out communication issues between exporters and VictoriaMetrics.

### **8.2.2 Access to the VMAAlert component**

VMAAlert has an HTTP interface where every fired alert is sent in a JSON response. This would help identify if a mistake was made in the VMAAlert or the VMAAlertmanager configuration.

### **8.2.3 Access to the VMAAlertmanager component**

VMAAlertmanager has a web interface where every information is available about fired alerts and some actions (e.g. silencing) can be taken. This would be extremely useful for debugging alert routing and notifications.

## **8.3 Improving Log Collection**

Log collection is very bare-bones right now and there is much room for improvement.

### **8.3.1 Configurations for Remote Promtails**

Installing promtails on remote servers is the lone duty of the remote sysadmin right now. User experience could be greatly improved by providing extra documentation and examples for our users.

### 8.3.2 Collecting Logs from Network Devices

At the time of writing no logs are being collected from network devices. An architecture similar to the SNMP exporters' could be a solution to this problem.

## 8.4 Advanced Log Analysis

Once the monitoring system has access to every metric and log it's going to have more than enough information to theoretically be able to detect complex problems and suggest possible solutions.

These issues are usually detected by our users themselves and then are reported in CESIG's support ticketing system. This is time consuming for both parties and can be a bad user experience.

Advanced log analysis could solve this and provide users with enough information to fix issues themselves preventing ticket creation.

Based on a crude analysis of tickets received during October 2022:

- 32 tickets were received in total
- 12 tickets could potentially be solved with this method
- 8 tickets would need the assistance of a member of CESIG but we could create alerts and/or provide useful information
- 11 tickets were requests that could be automated but not with the monitoring system
- 1 ticket could not be categorized

The second category is the one we are aiming for. A typical ticket of this category is about having no internet access. The dormitory's network denies access to unknown devices. Often the cause of this is MAC randomization which could easily be detected by analysing SNMP and/or network logs.

## 8.5 Service Discovery

With Advanced Config Generation (Section 8.1.4) adding new systems could be easy. But with service discovery this task could be (almost?) completely automated. There's nothing easier than doing nothing.

VictoriaMetrics supports many kinds of service discoveries. These methods should be explored and evaluated for a possible solution.

## 8.6 Other Possibilities

A few ideas that could be improved upon:

- Loki self-monitoring using Loki Canary
- Interactive messages for Mattermost
- Public vizualizations without authentication

## Chapter 9

# Conclusion

In this thesis I investigated previous implementations of monitoring systems in Computer Engineering Special Interest Group of Schönherz Zoltán Dormitory and researched other modern technologies that can handle this task. With CESIG 's members we decided to use VictoriaMetrics and Loki to achieve our goals. I designed an architecture around these softwares and implemented most parts of the new system. At the point of writing he proof of concept system was running for more than a month already. Members of CESIG started integrating their own servers with the help of the implemented Countinous Integration. The system proved to be successful so far.

During my work I learnt much about Kubernetes, monitoring systems, configuration management and Continous Integration.

There is still work to be done. In my opinion this solution has the potential to stay and server CESIG 's needs for a very long time after completion.

# Acknowledgements

I am grateful to my supervisor, Dr. Zoltán Zsóka for quickly accepting my last minute thesis application and for his continuous support since then.

I am extremely grateful to Computer Engineering Special Interest Group of Schönherz Zoltán Dormitory for trusting me with this important project's implementation and for helping me through my long journey on this university. Special thanks to the following members of CESIG :

- Miklós Tóth, the current Lead Sysadmin for sharing his ideas, expertise and valuable experiences on this topic
- Ádám Kiss, the next Monitoring System Administrator for supporting my work
- Tamás Kiss for sharing his related experiences, new ideas and for proofreading this thesis
- Adrián Robotka, Máté Turcsik and Balázs Ákos Kalmár for sharing their experiences on previous monitoring systems
- Kristóf Torma and Kornél Várkonyi for proofreading this thesis
- János Angyal, Józsué Majthényi-Wass, László Dorottya and László Radnai for keeping me company while writing this thesis

# Bibliography

- [1] Elasticsearch B.V. The elk stack: From the creators of elasticsearch (2022.12.09.). <https://www.elastic.co/what-is/elk-stack>.
- [2] Grafana Labs. Grafana: The open observability platform | grafana (2022.12.09.). <https://grafana.com/>, .
- [3] Grafana Labs. Grafana loki oss | log aggregation system (2022.12.09.). <https://grafana.com/oss/loki/>, .
- [4] Zoltán Janega. Felhasználói profiladat-források integrációja és ellenőrzött adat-átadása külső authorizációhoz, 05 2017.
- [5] Julien Menan. Monitoring at scale with victoria metrics (2022.09.06.). <https://tech.bedrockstreaming.com/2022/09/06/monitoring-at-scale-with-victoriametrics.html>.
- [6] Prometheus Authors. Prometheus - monitoring system & time series database (2022.12.09.). <https://prometheus.io/>.
- [7] The Kubernetes Authors. Kubernetes (2022.12.09.). <https://kubernetes.io/>.
- [8] VictoriaMetrics. Victorimetrics: Simple & reliable monitoring for everyone (2022.12.09.). <https://victoriametrics.com/>.

# Appendix

## A.1 Initial VictoriaMetrics values.yml

```
operator:
  enabled: true

serviceAccount:
  create: true

defaultRules:
  create: true
  rules:
    etcd: false

additionalVictoriaMetricsMap:

victoria-metrics-operator:
  operator:
    disable_prometheus_converter: false
    enable_converter_ownership: true
    psp_auto_creation_enabled: false
  rbac:
    pspEnabled: false

vmsingle:
  enabled: true
  # spec for VMSingle crd
  # https://github.com/VictoriaMetrics/operator/blob/master/docs/api.MD#
  vmsinglespec
  spec:
    retentionPeriod: "14"
    replicaCount: 1
    extraArgs: {}
    storage:
      resources:
        requests:
          storage: 20Gi
      storageClassName: openebs-zfspv

vmalert:
  enabled: false
```



```

vmagent:
  enabled: true
  additionalRemoteWrites: []
  spec:
    scrapeInterval: 25s
    externalLabels:
      cluster: beholder

# Grafana dependency chart configuration. For possible values refer to
# https://github.com/grafana/helm-charts/tree/main/charts/grafana#
# configuration
grafana:
  enabled: true
  sidecar:
    datasources:
      enabled: true
      createVMReplicasDatasources: false
      jsonData: {}
    dashboards:
      additionalDashboardLabels: {}
      additionalDashboardAnnotations: {}
      enabled: true
      multicluster: false

forceDeployDatasource: false

additionalDataSources: []

dashboardProviders:
  dashboardproviders.yaml:
    apiVersion: 1
    providers:
      - name: "default"
        orgId: 1
        folder: ""
        type: file
        disableDeletion: false
        editable: true
        options:
          path: /var/lib/grafana/dashboards/default

dashboards:
  default:
    nodeexporter:
      gnetId: 1860
      revision: 22
      datasource: VictoriaMetrics

defaultDashboardsEnabled: true

```

```

persistence:
  enabled: true
  type: statefulset
  size: 10Gi
  storageClassName: openebs-zfspv

ingress:
  enabled: true
  annotations:
    kubernetes.io/ingress.class: traefik
    kubernetes.io/tls-acme: "true"
    cert-manager.io/cluster-issuer: cert-issuer
    traefik.ingress.kubernetes.io/router.middlewares: monitoring-redirect-
https@kubernetescrd

  hosts:
    - beholder.sch.bme.hu
    - xn--ht8h.sch.bme.hu
  tls:
    - secretName: grafana-ingress-tls
      hosts:
        - beholder.sch.bme.hu
        - xn--ht8h.sch.bme.hu

vmServiceScrape:
  enabled: true
  spec: {}

kubeEtdcd:
  enabled: false # k3s does not have!

```

**Listing A.1.1:** victoria-metrics-k8s-stack.tex

## A.2 SNMP scraper helm chart

```
apiVersion: v2
name: snmp-exporter-target
description: Scrape configs for SNMP exporter
type: application
version: 0.0.0
appVersion: "yolo"
```

Listing A.2.1: ./Chart.yaml

```
{{- define "name" }}
{{- .Chart.Name }}-{{.Release.Name}}
{{- end }}
```

Listing A.2.2: ./templates/\_helpers.tpl

```
{{- $global := . }}
{{- range .Values.Devices }}

apiVersion: operator.victoriametrics.com/v1beta1
kind: VMStaticScrape
metadata:
  name: "{{ template "name" $global }}-{{ . }}"
spec:
  jobName: "{{ . }}"
  targetEndpoints:
    - targets:
        - "{{ template "name" $global }}"
      path: /snmp
      params:
        target: [ "{{ . }}" ]
        module: [ sch_core_net_mibs ]
{{- end }}
```

Listing A.2.3: ./templates/snmp-exporter.yaml

```

apiVersion: v1
kind: Service
metadata:
  name: "{{ template "name" . }}"
spec:
  ports:
    - port: 80
      targetPort: 9116
      name: exporter
      protocol: TCP
  type: ClusterIP

```

---

```

apiVersion: v1
kind: Endpoints
metadata:
  name: "{{ template "name" . }}"
subsets:
  - ports:
      - port: 9116
        name: exporter
        protocol: TCP
    addresses:
      {{- range .Values.NOCIPs }}
      - ip: {{ . }}
      {{- end }}

```

**Listing A.2.4:** ./templates/noc-snmp-lb.yaml

```

NOCIPs:
- 10.0.208.241
- 10.0.208.242

Devices:
- rtr-1

```

**Listing A.2.5:** ./values.yaml

## A.3 Alerts helm chart

```
apiVersion: v2
name: kszk-alerts
description: Alert configuration via helm for CI
type: application
version: 0.0.0
appVersion: "PESTILIENCE"
```

Listing A.3.1: ./Chart.yaml

```
apiVersion: operator.victoriametrics.com/v1beta1
kind: VMRule
metadata:
  name: kszk-alerts
  namespace: monitoring
spec:
  {{- .Values.spec | toYaml | trim | nindent 2 }}
```

Listing A.3.2: ./templates/kszk-alerts.yaml

```
apiVersion: operator.victoriametrics.com/v1beta1
kind: VMRule
spec:
  groups:
    - name: kszk-alerts
      rules: [ ]

# Very quick and very dirty CI solution for now.
# Feel free to modify anything under .spec but please don't alter others
# work without consultation.
# We are going to work on a much better solution in the near future.
```

Listing A.3.3: ./values.yaml